

ppOpen-HPC:

Open Source Infrastructure for Development and Execution of Large-Scale Scientific Applications on Post-Peta-Scale Supercomputers with Automatic Tuning (AT).

ppOpen-AT

ver. 1.0.0

Theory Manual

License

This software is an open source free software application. Permission is granted to copy, distribute and/or modify this software and document under the terms of The MIT license. A license file is included in the software archive.

This software is one of the results of the JST CREST ``ppOpen-HPC: Open Source Infrastructure for Development and Execution of Large-Scale Scientific Applications on Post-Peta-Scale Supercomputers with Automatic Tuning (AT)” project.

Change History

2012.Nov. Version 0.1 is released.

2013.Nov. Version 0.2 is released.

2016.Mar. Version 1.0.0 is released.

Changes in release 0.2.0

The following functions were added.

- ✓ Added a loop split and fusion function. See section 4, pp. 31 in this manual.
- ✓ Added a function for re-ordering of sentences. See section 4, pp. 33 in this manual.
- ✓ Fixed some bugs for adapting code in ppOpen-APPL/FDM and ppOpen-APPL/BEM version 0.1.0.

Introduction

ppOpen-AT is a domain specific language (set of directives) with features that reduce the workload for developers of libraries with auto-tuning features. Specifically, it is a scripting language (set of directives) whose purpose is to auto-generate code required for auto-tuning by placing annotations in the source program, in order to reduce the workload of library developers.

Although the concept of a library with auto-tuning features is applicable to a wide range of processes, the current specification is limited to a language specification specialized for parallel-calculation processing.

1. Target Languages

ppOpen-AT is a domain specific language (set of directives) designed to make developing parallel-computation libraries more efficient*. Consequently, it must have an interface from a language that is suited to numerical calculation and parallel processing. The purpose of *ppOpen-AT* is the generation of executable program code.

For these reasons, it is assumed that *ppOpen-AT* will be used to generate program code that will run in an environment in which Fortran90 or C is available as a language for numerical calculation, and Message Passing Interface (MPI) or/and OpenMP are available as a language for parallel processing.

* The language is not limited to library development. It can be used to optimize any user program.

1.2 Sequence of Library Development and Use

Fig. 1 illustrates the sequence of library development using *ppOpen-AT*, and the use of mathematics libraries with auto-tuning features by the user.

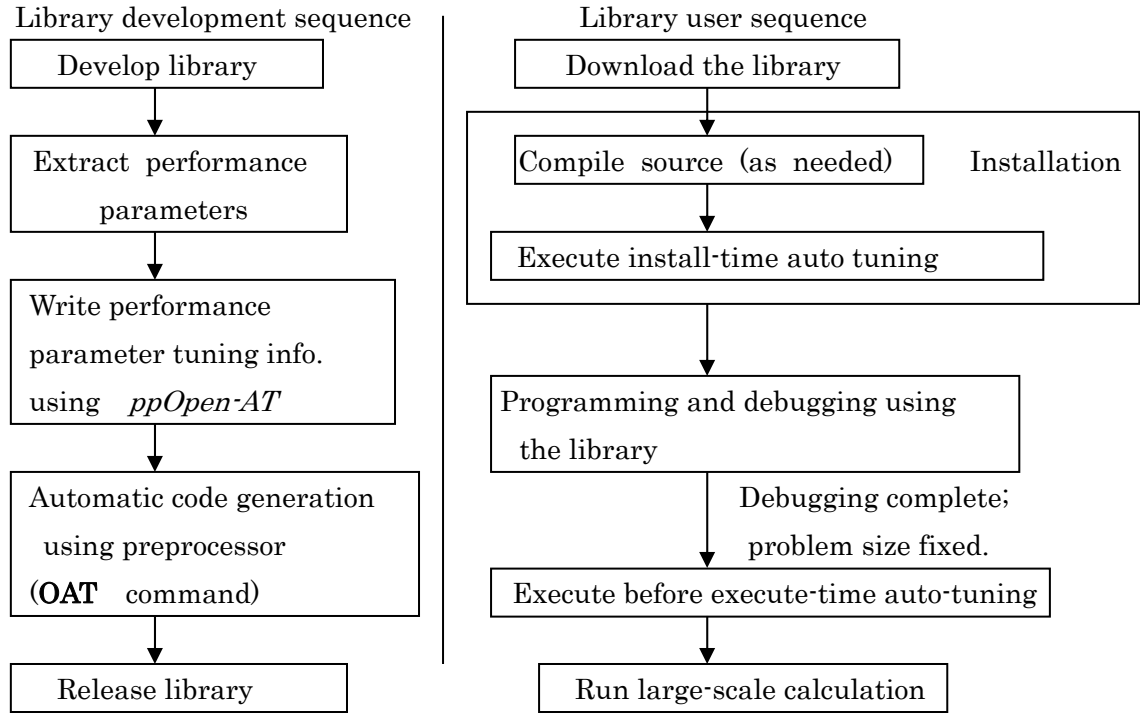


Figure 1. Sequence of Actions Required of Library Developers and Users Using *ppOpen-AT*.

1.3 Features of Programming with *ppOpen-AT*

ppOpen-AT was designed with the following four points in mind.

- (1) It is an auto-tuning specification function specialized for numerical calculation
Auto-tuning features optimized for numerical calculation are provided. They facilitate the creation of numerical calculation libraries with auto-tuning features.
- (2) It is a tuning specification method whereby the library developer inserts instructions into the program in the form of annotations.
The addition of auto-tuning features is performed in the form of annotations by the library developer. Consequently, this makes it easy to specify auto tuning, while not

interfering with the compilation of the original code.

- (3) A preprocessor that parses the library developer's annotations and automatically generates code is provided.

This feature parses the commands written by the library developer via annotations, and automatically generates code with added auto-tuning processing. This allows the library developer to manage the library source code itself, with an auto-tuning framework added on automatically. It is also possible to understand what processes have been added by viewing the automatically generated source code. This makes the processing highly transparent and worry-free.

- (4) The concept that two types of parameters affect performance: performance parameters (PP) and basic parameters (BP) is adopted.

The adoption of the concept of performance parameters and basic parameters improves the outlook for auto-tuning processing. It also makes it easy for library developers to communicate their intentions to the system.

1.4 The *ppOpen-AT* Software Architecture

The *ppOpen-AT* Application Programming Interface (API) consists of the following three elements:

- Specifiers defining auto-tuning targets, and subtype specifiers defining auto-tuning process details
- Environmental parameters defining auto-tuning information (system parameters)
- Run-time routines that supplement auto tuning

Fig. 2 illustrates the relationship between these elements.

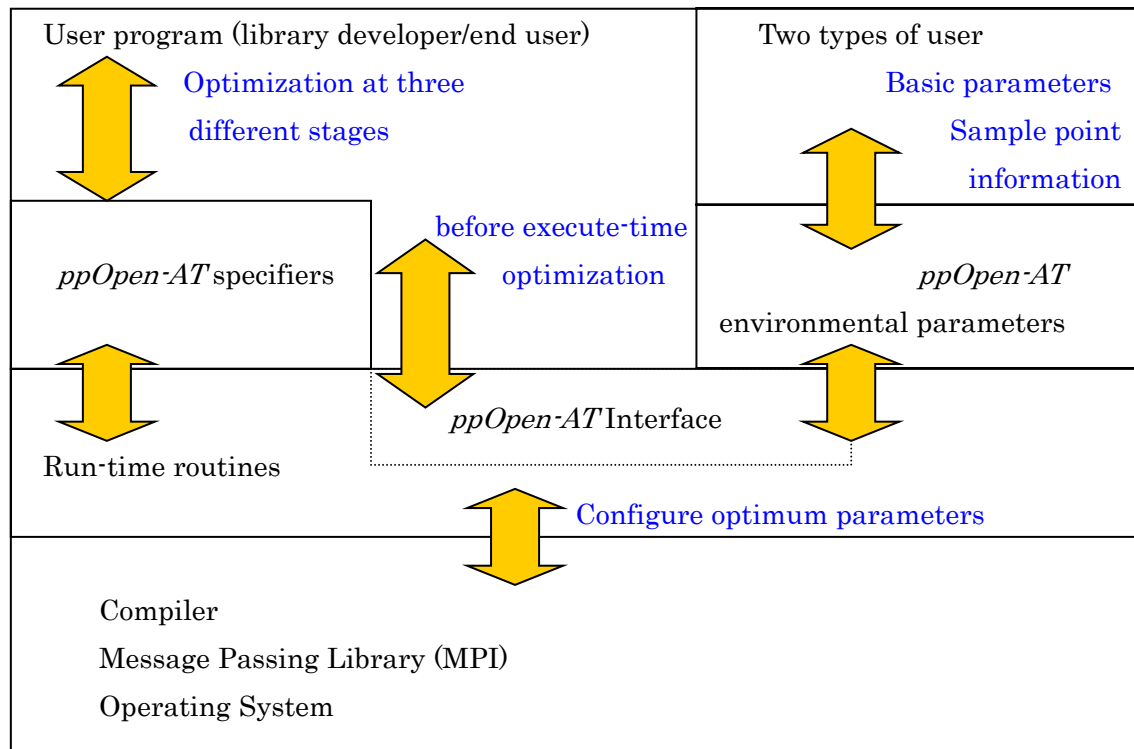


Figure 2. The *ppOpen-AT* Software Architecture

2. *Getting Started*

2.1 Auto-Tuning Features Envisioned by the System

The auto-tuning features provided by *ppOpen-AT* can be classified into the following three categories:

- **(1) Install-time auto-tuning**

This feature automatically tunes performance parameters that can be determined when the library is installed.

- **(2) Before execute-time auto-tuning**

This feature auto-tunes performance parameters after the end user has fixed the problem size and other basic parameters.

- **(3) Run-time auto-tuning**

This feature automatically tunes performance parameters that can be determined when the library is actually called from within a program.

These auto-tuning features are implemented as subroutines based on the Framework of Install-time, Before Execute-time and Run-time optimization feature (FIBER; patent application filed January 2003.) software architecture, an auto-tuning library architecture proposed by the ABCLibScript project, a precursor to the *ppOpen-AT* project.

Fig. 3 shows an overview of FIBER. With the exception of the auto-modeling routine shown in Fig. 3, *ppOpen-AT* is a scripting language (set of directives) developed to provide the three above-mentioned auto-tuning features.

Fig. 4 illustrates the parameter information hierarchy. In other words, parameters determined by the install-time auto-tuning routine can be referenced by the before execute-time auto tuning and run-time auto tuning routines. However, parameters determined by the before execute-time auto-tuning routine can only be referenced by the run-time auto-tuning routine. Finally, parameters determined by the run-time routine cannot be referenced by any other routine except the run-time routine[†].

[†] An exception, however, is in the FIBER feedback model, where it is possible for the before execute-time routine to reference parameters optimized by the run-time routine, and optimize them.

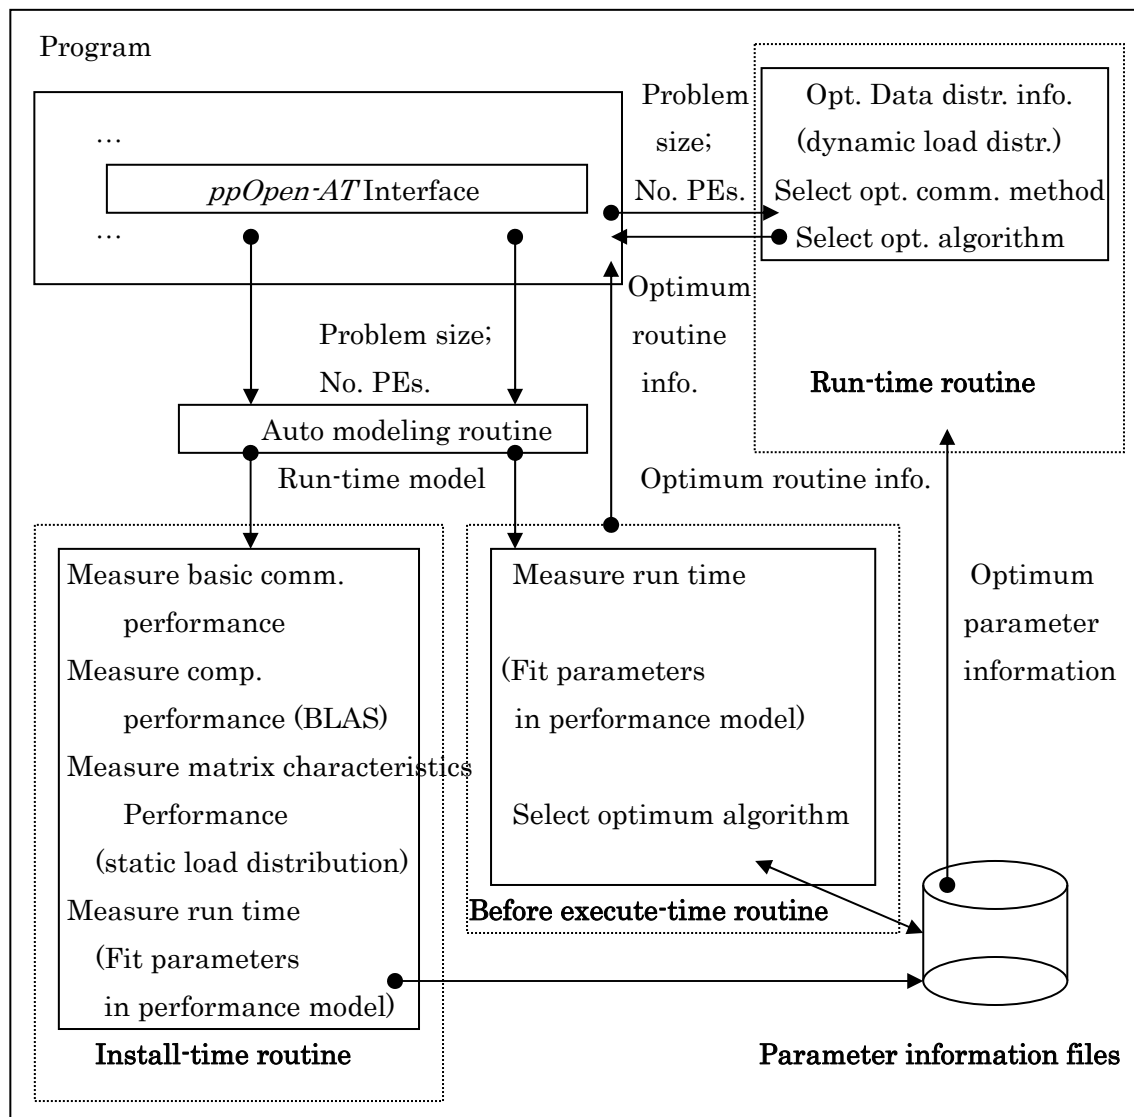


Figure 3. Structural Diagram of the Auto Tuning Library in the OAT Project
(FIBER Software Architecture)




Figure 4. Hierarchy of Parameter Information Referencing

2.2 Execution Priority for *ppOpen-AT*'s Auto Tuning

OAT executes auto-tuning in the following order.

- (1) Install-time auto tuning routine
- (2) Before execute-time auto tuning routine
- (3) Run-time auto tuning routine

Note that if the execution sequence is other than the above, an error code will be output, and auto-tuning will stop.

2.3 Parameter Types

FIBER tunes the following two types of parameter in libraries, subroutines, and in some programs developed by users (library developers and end users).

- **Basic parameters (*BP*)**

These parameters must be set when the end user uses the library. Some examples of BPs include matrix size and number of processors.

- **Performance parameters (*PP*)**

These parameters are not absolutely necessary for the end user to utilize the library, but do impact performance. One example is the number of unrolling levels. The user guarantees that the optimum values for performance information parameters can be found if the basic parameters are set.

ppOpen-AT is a scripting language (set of directives) developed to make it easy for library developers to specify the above-mentioned FIBER basic parameters and FIBER performance parameters. Consequently, the *ppOpen-AT* notation described below is a scripting language for declaring, defining, and specifying both performance parameters and basic parameters.

2.4 Notation

2.4.1 Overview

With *ppOpen-AT*, auto-tuning features are achieved by coding processes in the form of annotations in the source program. Specifically, lines beginning with

!OAT\$

are considered to be instructions to *ppOpen-AT*.

Briefly, the notational conventions are as follows:

```
!OAT$ <auto-tuning type> <feature name> [(target parameter)] region start
[ !OAT$ <feature details> [ sub region start ] ]
    Program targeted for processing
[ !OAT$ <feature details> [ sub region end] ]
!OAT$ <auto-tuning type> <feature name> [(target parameter)] region end
```

The <auto-tuning type> and <feature name> items above are called **specifiers**. The <feature details> items above are called **subtype specifiers**.

Finally, the processing surrounded by **!OAT\$ region start ... !OAT\$ region end** is called the **tuning region** or **AT region**.

2.4.2 Specifiers

The *ppOpen-AT* specifiers are described in detail below.

List of specifiers

<auto-tuning type>::= (**install** | **static** | **dynamic** | <formula>)

Install: Specifies install-time auto tuning

static: Specifies before execute-time auto tuning

dynamic: Specifies run-time auto tuning

<formula>::= Indicates a formula conforming to Fortran90 syntax

<feature name>::= (**define** | **variable** | **select** | **unroll**)

define: Specifies that this process sets a parameter

variable: Specifies that this is a variable parameter

select: Specifies that this process selects from multiple options

unroll: Specifies that the following process performs loop unrolling

2.4.3 Subtype Specifiers

The specifier <feature name> may require a further annotation. The subtype specifier is used to encode this. The following <feature details> subtype specifiers are available.

<feature details>::= (**name** | **parameter** | **select** | **according** |
 varied | **fitting** | **number** | **prepro** | **postpro**)

The <feature details> subtype specifiers that can be used depend on the specifier.

This is shown below.

Specifier	Subtype specifiers available
define	name, parameter, number, prepro, postpro
variable	name, parameter, varied, fitting, number, prepro, postpro
select	name, parameter, select, according, number, prepro, postpro
unroll	name, parameter, varied, fitting, number, prepro, postpro

Details about each subtype specifier are shown below.

List of subtype specifiers

- **name** <string>: Denotes the name of the tuning region
(available for all functions)
- **parameter** (<attribute specification> <parameter name>
[<attribute specification> <parameter name>,...])
Specify a parameter to output to a performance characteristics file, or input from a performance characteristics file

<attribute specification>::=[**in** | **out** | **bp**]

in: Input parameter (defined and referenced externally)

out: Output parameter (defined in this tuning region)

bp: Basic parameter

(Available for all functions)

- **select sub region** (**start** | **end**): Specifies that this is a selection procedure
(For the function-name select specification.(?This is ambiguous.))
- **according** (<conditional expression> | **estimated**): Specifies a selection procedure based on the standard specified below

<conditional expression>::=

[(**min**(<parameter name>) | **condition** (<condition>)) <connector>]

<connector>::=[**.and.** | **.or.**] <conditional expression>

estimated <mathematical expression>: Specifies that the optimum process should be selected and run, based on the user-defined cost (mathematical expression) associated with the selections.

(For the function-name select specification. (? This is ambiguous.))

- **Varied** (<parameter>[, <parameter>]) **from** X **to** Y
Specifies the range of variation of the specified parameter (from X to Y)
(For the (?) function-name variable/unroll specification.(?))

- **Fitting** <method> **sampler** <scope>: Specifies the method to use for inferring parameters

<method>::=

[**least-squares** <order> | **dspline** |

user-defined <mathematical expression> | **auto**]

least-squares: Specifies the parameter must be inferred by the least squares method via a polynomial expression.

<order> Sets the order of the polynomial expression.

dspline: Specifies the parameter must be inferred by discrete spline.

This method was developed by the Tanaka Laboratory, Kogakuin University.

user-defined: Infer using the least squares method, using the mathematical expression specified by the user.

auto: Allow the system to infer the parameter.

<scope>::= [<number> | **auto]** Specifies the scope required for parameter inference. Note that this can be omitted when **<method> = auto**

<number>: Specifies the concrete numerical value of the parameter

auto: Set the parameter's sampling interval automatically.

If the *fitting* subtype specifier is omitted, the optimum parameter is determined by measuring the entire range specified by the *varied* subtype specifier (i.e. exhaustive search).

(For the (?)function-name variable/unroll specification.(?))

- **number** **<number>**: Used to specify the order in which to process the tuning regions. If this is omitted, the regions are processed first to last. In the case of nested specifiers, a number can only be assigned to the outermost specifier (Available for all functions.)
- **prepro sub region (start | end)**: Specifies processing to apply before calling the tuning region (Available for all functions)
- **postpro sub region (start | end)**: Specifies processing to be applied after calling the tuning region (Available for all functions.)
- **debug** (**<parameter>**, [**<parameter>**, ...]): Specifies that the variable is to be shown in the debug display when the tuning region is executed

(Available for all functions.)

<parameter>::=[**bp | **pp** | **any**]**

bp: Display basic parameter information

pp: Display performance parameter information

Any: Display information about the specified parameter

Sample Program 1: Perform auto tuning upon installation to unroll a matrix product loop from 1 to 16 levels. The parameters are inferred by means of the least squares method, using a fifth-order polynomial equation. Additionally, set the sample points at 1-5, 8, and 16. Specify the performance parameter (number of unrolling levels) for debug display.

```
!OAT$ install unroll region start  
!OAT$ name MyMatMul  
!OAT$ varied (i,j) from 1 to 16  
!OAT$ fitting least-squares 5 sampled (1-5, 8, 16)  
!OAT$ debug (pp)  
do i=1, n  
  do j=1, n  
    do k=1,n  
       $A(i,j) = A(i,j) + B(i,k) * C(k,j)$   
    enddo  
  enddo  
enddo  
!OAT$ install unroll (i,j) region end
```

3. Implementing an Auto Tuning Library via *ppOpen-AT*

3.1 The Initialization Interface

ppOpen-AT provides an interface for executing auto tuning. The user can specify that auto tuning is desired by means of this interface.

Specifically, *ppOpen-AT* provides the following interface for auto tuning.

- **OAT_ATexec** (OAT_ATkinds, OAT_ATroutines)

The OAT_ATexec procedure performs the auto tuning specified by OAT_ATkinds within the auto-tuning region specified by OAT_ATroutines.

The OAT_ATkinds parameter is used to specify the type of auto-tuning to perform. One of the following four constants defined in the **OAT.h** header file can be specified.

OAT_INSTALL: Install-time auto tuning

OAT_STATIC: Before execute-time auto tuning

OAT_DYNAMIC: Run-time auto tuning

OAT_ALL: All auto-tuning

The OAT_ATroutines parameter specifies which tuning region to perform the processing on. This parameter can be declared by the user, using the type OAT_ATname defined in the header file OAT.h, or it can be specified using a global variable common-defined in OAT.h:

OAT_AllRoutines: For all routines

OAT_InstllRoutines: For install-time auto tuning routines

OAT_StaticRoutines: For before execute-time auto tuning routines

OAT_DynamicRoutines: For run-time auto tuning routines

The OAT_ATset subroutine substitutes auto-tuning processor information.

Note that during install-time auto tuning and before execute-time auto tuning, auto tuning is performed when OAT_ATexec is called. During run-time auto tuning, however, only execution is configured. Actual auto tuning is performed when the section in question is called.

Sample Usage

!OAT\$ call OAT_ATexec (OAT_INSTALL, OAT_InstallRoutines)

Performs install-time optimization.

- **OAT_ATset** (OAT_ATkinds, OAT_ATroutines)

The OAT_ATset procedure performs the procedures configured in OAT_ATroutines on the tuning region specified by the OAT_ATkind directive.

The OAT_ATkinds parameter is used to specify the type of auto-tuning to perform. One of the following four constants defined in the OAT.h header file can be specified.

OAT_INSTALL: Install-time auto tuning

OAT_STATIC: Before execute-time auto tuning

OAT_DYNAMIC: Run-time auto tuning

OAT_ALL: All auto-tuning

The name of the target tuning region is substituted in the OAT_ATroutines parameter.

Note that the following routine, OAT_ATdel, is used to delete specific routine names.

- **OAT_ATdel** (OAT_ATroutines, DelName)

The OAT_ATdel procedure deletes the tuning region with the name specified by DelName, from the parameter containing the tuning region name information specified by means of OAT_ATroutines.

In the DelName parameter, write the name of the tuning region to be deleted. Note that the tuning region name is (must be ??? must be でしょう.) written for every tuning region in the annotation format (see below for details).

Sample Usage

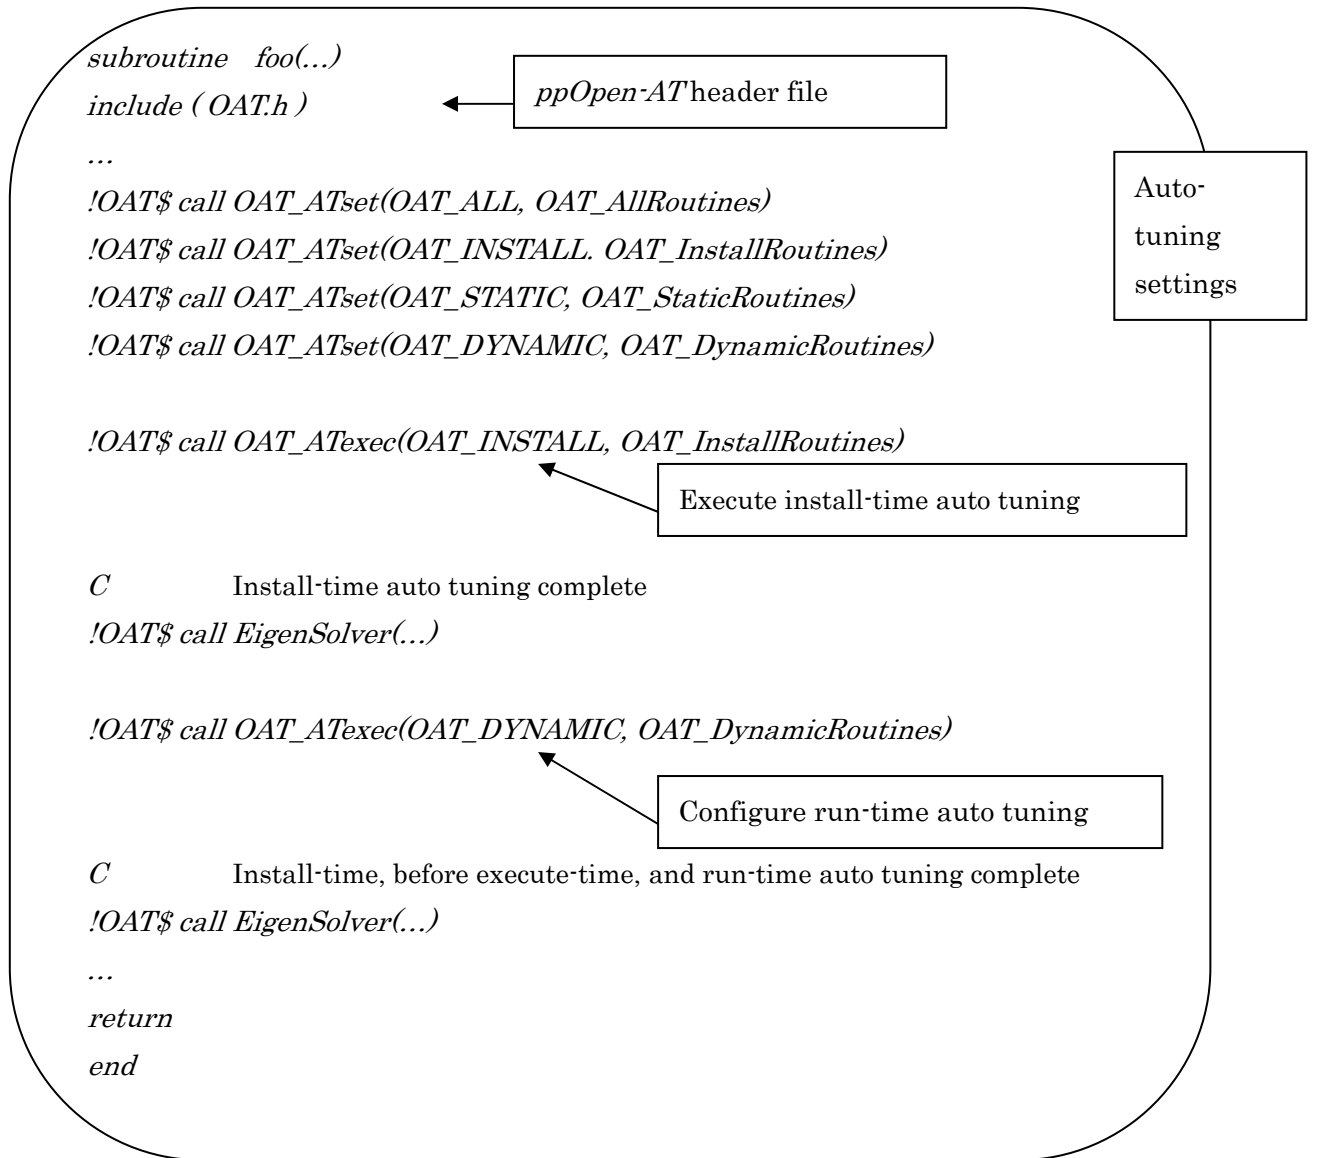
!OAT\$ call OAT_ATdel(OAT_InstallRoutines, "MyMatMul")

Deletes the MyMatMul tuning region from the candidates for install-time optimization.

3.2 Overview of Implementation Methods

This section describes each function, using concrete examples.

A library developer's subroutine, `EigenSolver`, can be written in the following sequence, using the *ppOpen-AT* interface in the library developer's subroutine, `foo`, which performs auto tuning.



Here, an end user is using the `EigenSolver` feature of the library provided by the library developer. Before execute-time auto tuning can be performed in the following sequence, using the *ppOpen-AT* interface in the library user's subroutine, `pooh`.

```
subroutine pooh(...)
```

```
include (OAT.h)
```

```
...
```

```
C Basic parameters specified by the library developer are constant
```

```
OAT_STARTTUNESIZE = 1234
```

```
OAT_ENDTUNESIZE=1234
```

```
call OAT_ATexec(OAT_STATIC, OAT_StaticRo(?)utines)
```

```
C Install-time and before execute-time auto tuning complete
```

```
call EigenSolver(...)
```

```
...
```

```
return
```

```
end
```

ppOpen-AT header file

Execute before execute-time auto tuning

3.2.1 Sample Use of Install-time Auto Tuning Routines

Install-time auto-tuning routines are generally meant to be run once after library installation. Consequently, by default, they are not executed the second and subsequent times they are specified. If you would like to process them again, you must initialize them by calling the OAT_ATInstallInit initialization interface below.

- **OAT_ATInstallInit** (OAT_InstallRoutines)

The OAT_ATInstallInit procedure undoes the tuning on the install-time tuning region specified in OAT_InstallRoutines.

Sample Usage

```
!OAT$ call OAT_ATInstallInit(OAT_InstallRoutines)
```

We assume that you want to run install-time auto tuning again.

In addition to the parameter inference performed upon installation as shown in **Sample Program 1**, install-time auto tuning routines are meant to measure those parameters that can be measured during installation. This specification is written using the following define function.

Sample Program 2:

Parameter determination using the define function in an install-time routine

```
!OAT$ install define ( CacheSize, CacheLine ) region start  
!OAT$ name SetCachePram(Spelling? SetCacheParam? See below.)  
!OAT$ parameter (out CacheSize, out CacheLine )  
  
...  
CacheSize =....  
CacheLine=....  
!OAT$ install define ( CacheSize, CacheLine ) region end
```

The install-time routine saves these parameter specification values in a **parameter information file**. The file name is OAT_InstallPram.dat (Spelling? Pram vs Param?), and it is stored in text format.

Example: Contents of OAT_InstallParam.dat (?See above.) parameter information file (SetCacheParam (?See above.)

```
(CacheSize 64)  
(CacheLine 8)  
)
```

3.2.2 Sample Use of Before execute-time Auto Tuning Routine

Before execute-time auto tuning determines parameters based on detailed information before the end user executes the library. Before execute-time auto tuning tunes parameters on the condition that the end user guarantees the values of the basic parameter described in section 2.4.

The user specifies the parameters to be provided before library execution in the OAT_StaticParamDef.dat parameter information file. Basic parameters can be configured via substitution statements in the program. Here, **basic parameters** are required for install-time and before execute-time auto tuning. By default, they are referred to internally using parameter names specified by the system.

The default basic-information parameter settings are shown below.

List of Default Basic Parameters

<Default Basic Parameters>::=

(OAT_NUMPROCS | OAT_STARTTUNESIZE | OAT_ENDTUNESIZE | OAT_SAMPDIST)

OAT_NUMPROCS: Integer. Indicates the number of processors to use.

OAT_STARTTUNESIZE: Integer. The starting problem size for auto tuning with regard to the default basic parameters.

OAT_ENDTUNESIZE: Integer. The ending problem size for auto tuning with regard to the default basic parameters.

OAT_SAMPDIST: Integer. The size of the increment used to increase the problem size for auto tuning with regard to the default basic parameters.

Note that before execute-time auto tuning will not run if the basic parameters are not set. Additionally, install-time auto tuning will not run unless OAT_NUMPROCS, OAT_(?)STARTTUNESIZE, OAT_ENDTUNESIZE, and OAT_SAMPDIST are set.

Library developers wishing to set new basic parameters should use the OAT_BPset and OAT_BPsetName procedures below.

- **OAT_BPset(BPvalName)**

The OAT_BPset procedure makes the basic parameter name specified by the parameter BPvalName into a new basic parameter.

Sample Usage

!OAT\$ call OAT_BPset ("nprocs")

Makes the parameter nprocs into a basic parameter.

- **OAT_BPsetName**(Kind, BPvalName, Name)

The OAT_BPsetName procedure sets the name of the information parameter relating to the basic parameter of the type specified by the parameter Kind to the name specified by the parameter Name, for the basic parameter specified in the parameter BPvalName.

Here, the parameter Kind is as follows.

Kind ::= [**STARTTUNESIZE** | **ENDTUNESIZE** | **SAMPDIST**]

STAT(<=R?)TTUNESIZE: Sample start point information relating to
basic parameter BPvalName

ENDTUNESIZE: Sample end point information relating to
basic parameter BPvalName

SAMPDIST: Sample point interval information relating to
basic parameter BPvalName

Sample Usage

```
!OAT$ call OAT_BPsetName("STARTTUNESIZE", "nprocs",  
!OAT$ & "OAT_NprocsStartSize")
```

Sets the sample start point parameter name for auto tuning for basic parameter nprocs to OAT_NprocsStartSize.

- **OAT_BPsetCDF**(BPvalName, CFKind)

The OAT_BPsetCDF (CDF vs CFD? See below.) procedure sets the method to be used for inferring non-sample points relating to the basic parameter name specified in the parameter BPvalName to the type of cost definition function specified by the parameter CFD(?)Kind. Here, the parameter CFD(?)Kind is as follows.

CFD(?)Kind ::= [**least-squares** <order> | **user-defined** <mathematical expression> | **auto**]

least-squares <order>: Specifies the parameter must be inferred by the least squares method via a polynomial expression. <order> Sets the order of the

polynomial expression.

user-defined <mathematical expression>: Infer using the least squares method,
using the mathematical expression specified by the user

auto: Allow the system to infer the parameter.

Note that, by default, the cost definition function specified for the tuning region in question is used as-is.

Sample Usage

```
!OAT$ call OAT_BPsetCFD(?)("nprocs", "least-squares 5")
```

Infer parameters for basic parameter nprocs using the least squares method, as a fifth-order polynomial equation.

Sample Program 3: Specify the basic-parameter sample points when performing before execute-time tuning for 1,024, 2,048, and 3,072 dimensional parameters using 4 processors.

```
!OAT$ OAT_TUNESTATIC = 1  
!OAT$ OAT_NUMPROCS = 4  
!OAT$ OAT_STARTTUNESIZE = 1024  
!OAT$ OAT_ENDTUNESIZE = 3072  
!OAT$ OAT_SAMPDIST = 1024  
!OAT$ call OAT_ATexec(OAT_STATIC, OAT_StaticRoutines)
```

Another way of specifying this is to write the following in the OAT_StaticParam(?)Def.dat file.

```
(BasicParam  
  (OAT_TUNESTATIC 1)  
  (OAT_NUMPROCS 4)  
  (OAT_STARTTUNESIZE 1024)  
  (OAT_ENDTUNESIZE 3072)  
  (OAT_SAMPDIST 1024)  
)
```

It is expected that the unroll feature will be used in before execute-time routines. Below is a sample of this usage.

Sample Program 4a:

This before execute-time routine determines the optimum parameter, unrolling the loop in the program below to 16 levels. Parameter measurement is performed at each of levels 1 to 16 (i.e. exhaustive search). Note, however, that it is assumed that the sample points for the basic parameter have been set as in **Sample Program 3**. Since the only loop entry variable is n , parameter n is interpreted as the default basic parameter.

```
!OAT$ static unroll (i, j) region start
!OAT$ name MyMatMul
!OAT$ varied (i, j) from 1 to 16
do i=1, n
  do j=1, n
    do k=1, n
       $A(i, j) = A(i, j) + B(i, k) * C(k, j)$ 
    enddo
  enddo
enddo
!OAT$ static unroll (i, j) region end
```

In the case of this example, the optimum output parameters (i.e. the parameters for which processing is fastest) are the values i, j .

The before execute-time auto tuning routine writes this optimized parameter to the OAT_StaticParam(?).dat parameter information file.

Example: Contents of OAT_StaticParam(?).dat parameter information file:

```
(MyMatMul
(OAT_NUMPROCS 4)
(OAT_SAMPDIST 1024)
(OAT_PROBSIZE 1024
(MyMatMul_I 4)
(MyMatMul_J 8) )
```

```

(OAT_PROBSIZE  2048
  (MyMatMul_I   4)
  (MyMatMul_J   9) )
(OAT_PROBSIZE  3072
  (MyMatMul_I   5)
  (MyMatMul_J  10) )
)

```

Before execute-time auto-tuning routine parameter determination can be performed only by using parameters specified in the program, or (???) by calling the parameters of the install-time auto-tuning routine (the data in the parameter information file). **Sample Program 5** shows an example of this.

Sample Program 4b:

This before execute-time routine determines the optimum parameter, unrolling the loop in the program below to 16 levels. Parameter measurement is performed at each of levels 1 to 16 (i.e. exhaustive search). Note, however, that it is assumed that the basic information parameters have been specified as in **Sample Program 3**.

In the example below, because n is not the only loop-termination variable (the variable $nprocs$ also exists), it is not known whether the variable n or $nprocs$ is the default basic parameter. Consequently, the user must specify the basic parameter via the parameter subtype specifier.

```

!OAT$ static unroll (i,j) region start
!OAT$ name MyMatMul
!OAT$ parameter(bp n)
!OAT$ varied (i,j) from 1 to 16
do i=1, n/nprocs
  do j=1, n
    do k=1, n
       $A(i,j) = A(i,j) + B(i,k) * C(k,j)$ 
    enddo
  enddo
enddo
!OAT$ static unroll (i,j) region end

```


Sample Program 4c:

This before execute-time routine determines the optimum parameter, unrolling the loop in the program below to 16 levels. Parameter measurement is performed at each of levels 1 to 16 (i.e. exhaustive search). Note, however, that it is assumed that the basic information parameters have been specified as in **Sample Program 3**. In the example below, both n and $nprocs$ are specified as basic parameters.

```
!OAT$ call OAT_BPsetVal("nprocs")
!OAT$ call OAT_BPsetName(STARTTUNESIZE, "nprocs",
!OAT$ &          "OAT_NprocsStartSize")
!OAT$ call OAT_BPsetName(ENDTUNESIZE, "nprocs",
!OAT$ &          "OAT_NprocsEndSize")
!OAT$ call OAT_BPsetName(SAMPDIST, "nprocs",
!OAT$ &          "OAT_NprocsSampDist")
!OAT$ OAT_NprocsStartSize = 1
!OAT$ OAT_NprocsEndSize = 8
!OAT$ OAT_NprocsSampDist = 1
!OAT$ static unroll (i, j) region start
!OAT$ name MyMatMu
!OAT$ parameter(bp n, bp nprocs)
!OAT$ varied (i, j) from 1 to 16
do i=1, n/nprocs
  do j=1, n
    do k=1, n
      
$$A(i, j) = A(i, j) + B(i, k) * C(k, j)$$

    enddo
  enddo
enddo
!OAT$ static unroll (i, j) region end
```

Sample Program 5:

Determine the optimum implementation method referring to the parameter CacheSize of the install-time auto-tuning routine, and using the basic parameter information for problem size and number of processors set by the end user before execution. Here, the selection is made automatically, using a given standard (in this case, the run-time estimate provided by the user)[‡].

```
!OAT$ static select region start
!OAT$ name ATfromCacheSize
!OAT$ parameter (in CacheSize, in OAT_PROBSIZE,
!OAT$ & in OAT_NUMPROC)
!OAT$ select sub region start
!OAT$ according estimated
!OAT$ & 2.0d0*CacheSize*OAT_PROBSIZE*OAT_PROBSIZE
!OAT$ & / (3.0d0*OAT_NUMPROC)
Target process 1
!OAT$ select sub region end
!OAT$ select sub region start
!OAT$ according estimated 4.0d0*CacheSize*OAT_PROBSIZE
!OAT$ & *dlog(OAT_PROBSIZE) / (2.0d0*OAT_NUMPROC)
Target process 2
!OAT$ select sub region end
!OAT$ static select region end
```

3.2.3 Sample Use of a Run-time Auto Tuning Routine

Run-time auto tuning routines determine performance parameters based on information that can be obtained at run-time. When determining these performance parameters, it is possible to reference parameters specified in the program, parameters determined by an install-time auto-tuning routine, and parameters specified by a before execute-time auto-tuning routine.

[‡] In this case, the selection is based on execution time. Note that execution time is generally not the only standard for AT region selection. For example, subroutine cost is also a possible selection standard. This type of selection standard can be encoded using the select specifier.

Below is a description of the select function, the use of which can be expected in run-time routines.

Sample Program 6:

Select the optimum process based on the parameters *eps* and *iter* defined in the program. Specifically, set condition *iter* to less than 5, in order to make *eps* the minimum value.

```
!OAT$ dynamic select (eps, iter) region start
!OAT$ name PricondSelect
!OAT$ parameter (in eps, in iter)
!OAT$ according min (eps) .and. condition (iter < 5)
!OAT$ select sub region start
    Target process 1
    eps=...
!OAT$ select sub region end
!OAT$ select sub region start
    Target process 2
    eps=...
!OAT$ select sub region end
!OAT$ dynamic select (eps, iter) region end
```

● OAT_DynPefThis(Name)

The OAT_DynPefThis procedure specifies that the optimization of the tuning region name performed by the run-time auto-tuning routine specified in the parameter Name is to be performed at the location written in this procedure. Note that when using this procedure, the execution of ((auto) tuning for??)the tuning region n(N?)ame is performed using optimized parameters. In other words, no parameter tuning is performed by the tuning region specified by Name.

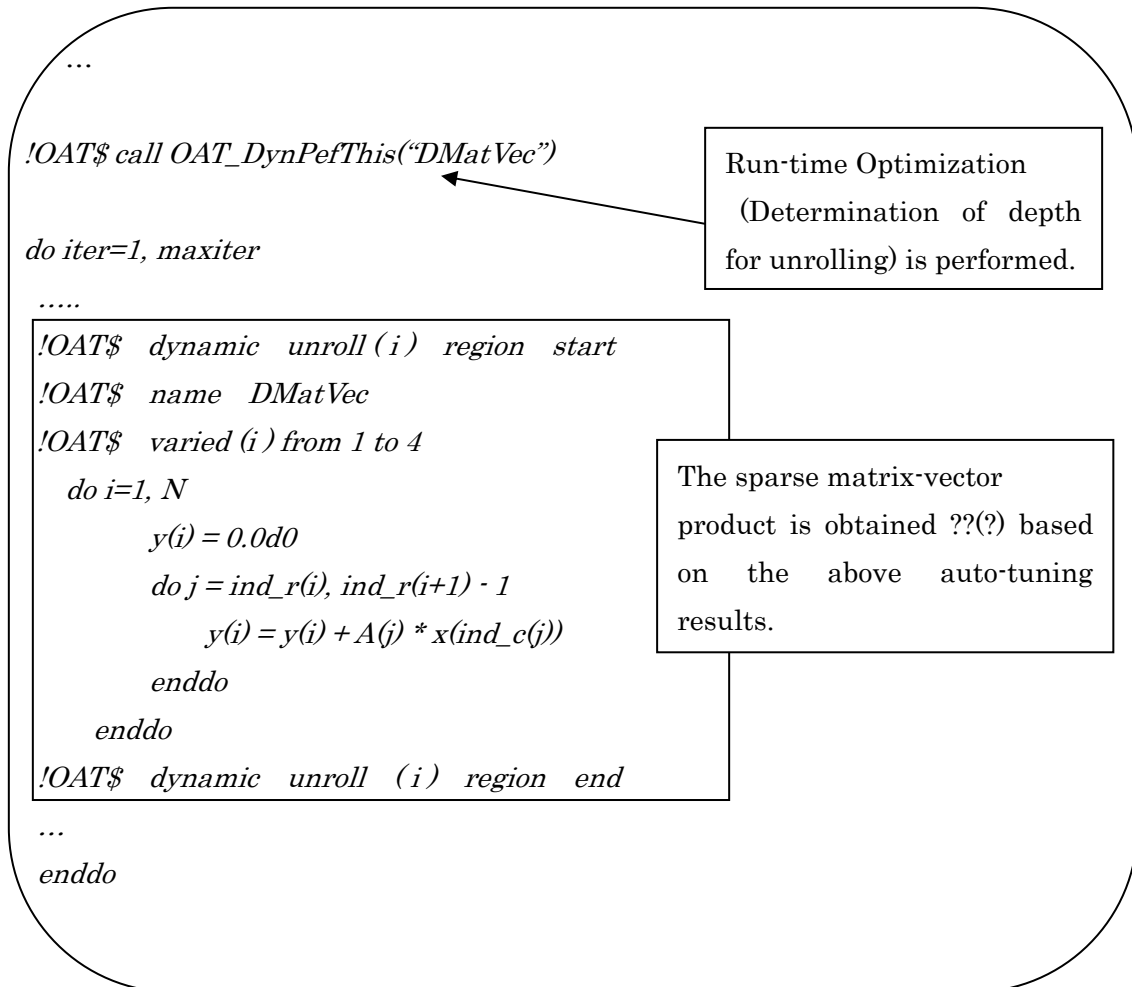
Sample Usage

```
!OAT$ call OAT_DynPefThis ( "DMatVec" )
```

Perform the auto tuning of the auto tuning region DMatVec here.

Sample Program 7:

Regarding the determination of the optimum number of loop-unrolling levels, re-use the parameters optimized by run-time optimization, and execute the process optimized in the AT region in question.



3.3 Sample Use of the Auto Tuning Code Generation Command

To generate actual parallel-computing library code with auto-tuning features (Fortran90 + MPI), do the following:

```
>OATCodeGen test.f
```

Here we assume that test.f is a program including ppOpen-AT directives.

After executing the command above, an OAT directory is created beneath the current directory, in which the following files are created:

./OAT/OAT_test.f:

Fortran90 + MPI source code with embedded auto-tuning code in test.f

./OAT/OAT_InstallRoutines.f:

Subroutines extracted from test.f that perform install-time auto tuning

./OAT/OAT_StaticRoutines.f:

Subroutines extracted from test.f that perform (?) before execute-time auto tuning

./OAT/OAT_DynamicRoutines.f:

Subroutines extracted from test.f that perform (?) run-time auto tuning

./OAT/OAT_ControlRoutines.f:

Subroutines that control auto-tuning code

If these files already exist when the preprocessor starts, non-overlapping source code and subroutines are added to the above files.

3.3.1 Run-time Options

OATCodeGen can specify the following run-time options.

—debug

OFF: Do not generate debugging code

ON: Generate code debugging code at debug level x, specified by the OAT_PRINT parameter.

—visualization

OFF: Do not output an auto-tuning trace file (default value)

ON: Output an auto-tuning trace file, and use it to create a visualization

The name of the auto-tuning trace file created when **—visualization** is set to ON is:

- OATATlog.dat

Sample Usage

```
> OAT -debug ON -visualization ON test.f
```

Specify the generation of debugging code, and the output of an auto-tuning trace file of the trace of auto-tuning mode for program test.f, containing *ppOpen-AT* specifiers. This auto-tuning trace file can be used to dynamically view tuning progress in a visualizer (*VizOAT*).

4. *Extended Functions*

4.1 Overview

In versions after version 0.2, or in other later versions, new functions for loop transformation are available. The main functions are loop split and loop fusion with data dependences. In addition to that, re-ordering of sentences in a loop is also available. In this session, we explain two key examples for the new extended functions to the original functions of ppOpen-AT.

4.2 Functions for loop split and fusion with data dependences

The ppOpen-AT functions for loop split and loop fusion are explained with the following Sample Program 8.

Sample Program 8:

The following is an example of how to adapt the loop split and fusion functions to the heaviest kernel in ppOpen-APPL/FDM. There is a flow dependency between the definition of QG and several uses of QG. Hence in general, it is difficult to perform loop splitting using compilers.

The new directive for specifying loop split and loop fusion is “*!oat\$ install LoopFusionSplit region start*” ~ “*!oat\$ install LoopFusionSplit region end*”.

To define a loop split point, “*!oat\$ SplitPoint (K, J, I)*” can be used. The target loops, including loop induction variables, for splitting can be defined with the directive. In this case, loops for K, J, and I can be adapted for the loop splitting operation.

The re-computation sentences when a loop split is done can specified with “*!oat\$ SplitPointCopyDef region start*” ~ “*!oat\$ SplitPointCopyDef region end*”. Points to which to copy the sentences that are defined by “*!oat\$ SplitPointCopyDef region start*” ~ “*!oat\$ SplitPointCopyDef region end*” can be defined with “*!oat\$ SplitPointCopyInsert*”.

```

!oat$ install LoopFusionSplit region start
!$omp parallel do private (k, j, i, STMP1, STMP2, STMP3, STMP4, RL, RM, RM2,
    RMAXY, RMAXZ, RMAYZ, RLTHETA, QG )
DO K = 1, NZ
DO J = 1, NY
DO I = 1, NX
    RL = LAM (I,J,K);  RM = RIG (I,J,K);  RM2 = RM + RM
    RLTHETA = (DXVX(I,J,K)+DYVY(I,J,K)+DZVZ(I,J,K))*RL
!oat$ SplitPointCopyDef region start
    QG = ABSX(I)*ABSY(J)*ABSZ(K)*Q(I,J,K)
!oat$ SplitPointCopyDef region end
    SXX (I,J,K) = ( SXX (I,J,K) + (RLTHETA + RM2*DXVX(I,J,K))*DT ) * QG
    SYX (I,J,K) = ( SYX (I,J,K) + (RLTHETA + RM2*DYVY(I,J,K))*DT ) * QG
    SZZ (I,J,K) = ( SZZ (I,J,K) + (RLTHETA + RM2*DZVZ(I,J,K))*DT ) * QG
!oat$ SplitPoint (K, J, I)

    STMP1 = 1.0/RIG(I,J,K);  STMP2 = 1.0/RIG(I+1,J,K);  STMP4 = 1.0/RIG(I,J,K+1)
    STMP3 = STMP1 + STMP2
    RMAXY = 4.0/(STMP3 + 1.0/RIG(I,J+1,K) + 1.0/RIG(I+1,J+1,K))
    RMAXZ = 4.0/(STMP3 + STMP4 + 1.0/RIG(I+1,J,K+1))
    RMAYZ = 4.0/(STMP3 + STMP4 + 1.0/RIG(I,J+1,K+1))
!oat$ SplitPointCopyInsert
    SXY (I,J,K) = ( SXY (I,J,K) + (RMAXY*(DXVY(I,J,K)+DYVX(I,J,K)))*DT ) * QG
    SXZ (I,J,K) = ( SXZ (I,J,K) + (RMAXZ*(DXVZ(I,J,K)+DZVX(I,J,K)))*DT ) * QG
    SYZ (I,J,K) = ( SYZ (I,J,K) + (RMAYZ*(DYVZ(I,J,K)+DZVY(I,J,K)))*DT ) * QG
END DO
END DO
END DO
!$omp end parallel do
!oat$ install LoopFusionSplit region end

```

By adapting the preprocessor of ppOpen-AT, we can obtain candidates for code that provide all combinations by adapting loop split and loop fusion by using the above directives. In the case of sample 8, we can obtain the following 8 candidates:

- #1 [Baseline] : Original three-nested loop.
- #2 [Spilt] : Loop split for the k-loop (separated into two three-nested loops).
- #3 [Split] : Loop split for the j-loop.
- #4 [Split] : Loop split for the i-loop.
- #5 [Fusion] : Loop fusion for the k-loop and j-loop (a two-nested loop).
- #6 [Split and Fusion] : Loop fusions for the k-loop and j-loop for the loops in #2.
- #7 [Fusion] : Loop fusions for the k-loop, j-loop, and i-loop (loop collapse).
- #8 [Split and Fusion] : Loop fusions for the k-loop, j-loop, and i-loop for the loops in #2 (the loop collapses for the two separated loops).

4.3 Function for re-ordering of sentences in a loop

The ppOpen-AT function for re-ordering of sentences is explained in the following Sample Program 9.

Sample Program 9:

The following is an example of how to adapt re-ordering of sentences in loops. In this example, the functions of loop split and fusion to one of the heavy kernels in ppOpen-APPL/FVM is applied. In this case, loop fusion is specified, but you can also specify other loop transformations, such as a loop split only.

The target sentences for the re-ordering of sentences can be specified with “!OAT\$ RotationOrder sub region start”~ “!OAT\$ RotationOrder sub region end”. This specification is needed for two regions. Each sentence that is specified the directive is re-ordered.


```

!OAT$ install LoopFusion region start
do k = NZ00, NZ01
  do j = NY00, NY01
    do i = NX00, NX01
!OAT$ RotationOrder sub region start
      ROX = 2.0_PN/( DEN(I,J,K) + DEN(I+1,J,K) )
      ROY = 2.0_PN/( DEN(I,J,K) + DEN(I,J+1,K) )
      ROZ = 2.0_PN/( DEN(I,J,K) + DEN(I,J,K+1) )
!OAT$ RotationOrder sub region end
!OAT$ RotationOrder sub region start
      VX(I,J,K) =
        VX(I,J,K) + ( DXSXX(I,J,K)+DYSXY(I,J,K)+DZSXZ(I,J,K) )*ROX*DT
      VY(I,J,K) =
        VY(I,J,K) + ( DXSXY(I,J,K)+DYSYY(I,J,K)+DZSYZ(I,J,K) )*ROY*DT
      VZ(I,J,K) =
        VZ(I,J,K) + ( DXSXZ(I,J,K)+DYSYZ(I,J,K)+DZSZZ(I,J,K) )*ROZ*DT
!OAT$ RotationOrder sub region end
    end do
  end do
end do
!OAT$ install LoopFusion region end

```

One of code segments generated for the original 3-nested loop is as follows.

```

do k = NZ00, NZ01
  do j = NY00, NY01
    do i = NX00, NX01
      ROX = 2.0_PN/( DEN(I,J,K) + DEN(I+1,J,K) )
      VX(I,J,K) = VX(I,J,K) +DXSXX(I,J,K)+DYSXY(I,J,K)+DZSXZ(I,J,K) )*ROX*DT
      ROY = 2.0_PN/( DEN(I,J,K) + DEN(I,J+1,K) )
      VY(I,J,K) = VY(I,J,K) + ( DXSXY(I,J,K)+DYSYY(I,J,K)+DZSYZ(I,J,K) )*ROY*DT
      ROZ = 2.0_PN/( DEN(I,J,K) + DEN(I,J,K+1) )
      VZ(I,J,K) = VZ(I,J,K) + ( DXSXZ(I,J,K)+DYSYZ(I,J,K)+DZSZZ(I,J,K) )*ROZ*DT
    end do
  end do
end do

```

end do
end do

5. *ppOpen-AT Internal Specifications*

5.1 System Parameters

Below is a list of system parameters kept by *ppOpen-AT*. Note that these parameters are reserved words, and cannot be defined by the user.

List of System Parameters (Reserved Words)

Tuning Type Specifiers

<p>OAT_ALL: An integer type with value 0. Indicates all tuning types (install-time, before execute-time, and run-time)</p> <p>OAT_INSTALL: An integer type with value 1. Indicates install-time optimization.</p> <p>OAT_STATIC: An integer type with value 2. Indicates before execute-time optimization.</p> <p>OAT_DYNAMIC: An integer type with value 3. Indicates run-time optimization.</p>

Tuning Region Name Storage

<p>OAT_AllRoutines: A string type. All tuning region names are stored.</p> <p>OAT_InstallRoutines: A string type. The names of tuning regions for which install-time auto tuning is to be performed are stored.</p> <p>OAT_StaticRoutines: A string type. The names of tuning regions for which before execute-time auto tuning is to be performed are stored.</p> <p>OAT_DynamicRoutines: A string type. The names of tuning regions for which run-time auto tuning is to be performed are stored.</p>

Default Basic Parameters

<p>OAT_NUMPROCS: Integer. Holds the number of processors to be used.</p> <p>OAT_STARTTUNESIZE: Integer. Holds the starting sample point value for the default basic parameters.</p> <p>OAT_ENDTUNESIZE: Integer. Holds the ending sample point value for the default basic parameters.</p> <p>OAT_SAMPDIST: Integer. Holds the sample point interval (increment) value for the default basic parameters.</p>
--

System Control

OAT_TUNESTATIC: Boolean. Specifies whether to execute before execute-time auto tuning.

.true.: Execute. **.false.:** Do not execute.

OAT_TUNEDYNAMIC: Boolean. Specifies whether to execute run-time auto tuning.

.true.: Execute. **.false.:** Do not execute.

OAT_DEBUG: Integer. Specifies the debug print level.

0: None. **Value x of 1 or greater:** Debug printing at level x.

4.2 Input and Output Files

This section describes the input and output files (parameter information files) handled by ppOpen-AT. There are two types of I/O file: Files automatically generated by ppOpen-AT (**system specification files**) and files specified by the user for debugging and the like (**user specification files**).

System specification files:

OAT_InstallParamX.dat: For install-time auto tuning routine parameter output

OAT_StaticParamX.dat: For before execute-time auto tuning routine parameter output

User specification files

OAT_InstallParamDefX.dat: For install-time auto tuning routine parameter specification

OAT_StaticParamDefX.dat: For before execute-time auto tuning routine parameter specification

OAT_DynamicParamDefX.dat: Input file for before execute-time auto tuning routine parameter specification

Note: *X* holds the name of the AT region in question.

5.2.1 Input Files

Input files consist of both user specification files and system specification files.

User specification files

OAT_InstallParamDefX.dat: For install-time auto tuning routine parameter specification

OAT_StaticParamDefX.dat: For before execute-time auto tuning routine
parameter specification

OAT_DynamicParamDefX.dat: For run-time auto tuning routine
parameter specification

System specification files:

install-time routines: None

Before execute-time routines:

OAT_InstallParamX.dat

Run-time routines:

OAT_StaticParamX.dat,

OAT_DynamicParamX.dat

5.2.2 Output Files

Output files consist of system specification files only.

System specification files:

install-time routines:

OAT_InstallParamX.dat

Before execute-time routines:

OAT_StaticParamX.dat

Run-time routines:

OAT_DynamicParamX.dat

5.2.3 Inupt/Output File Format

The format of user specification files and system specification files is as follows.

```
<format>::=
(<name>
  (<key> <value>)
  [(<key> > <value>)]
  ...
)
[<format>];
<key>::= (<format> | parameter name) ;
<value>::=[parameter value];
```

<name> specifies a tuning region name or basic parameter name. To specify a basic parameter, write BasicParam.

5.3 Collisions with Parameters in User Specification Files

When auto tuning is performed on a parameter specified in a user specification file (when an attempt is made to determine the parameter), this is called a **parameter collision**.

When there is a parameter collision, auto tuning halts, and the user-specified parameter is forcibly set.

When the system detects a parameter collision, it assumes that the user wants to halt the auto-tuning feature in order to debug or the like. To put it the other way around, the user can perform debugging by defining parameter information in a user specification file.

5.4 Nesting of Statements

This section defines specifier nesting.

5.4.1 Nesting Availability and Depth

The type of nesting available in auto-tuning is defined in Table 1.

Table 1. Availability of Nesting by Auto Tuning Type			
Nesting part (superior part)	Nested part (subordinate part)		
	install	static	dynamic
install	yes	no	no
static	yes	yes	no
dynamic	yes	yes	yes

Table 2 defines combinations of features that can be nested.

Table 2. Nesting Availability by Feature

Nesting part (superior part)	Nested part (subordinate part)			
	define	variable	select	unroll
define	yes	yes	yes	yes
variable	yes	yes	yes	yes
select	yes	yes	yes	yes
unroll	no	no	no	no

The maximum nesting depth (how far down elements can be nested) is currently as shown below. In other words, the maximum nesting depth is three.

Nesting depth = 3 or fewer

5.4.2 Parameter Search Order (extended feature)

The method used for searching for nested parameters is determined by the method specified by the outermost tuning region. The parameter search method can be annotated as follows.

subtype specifier <parameter search method>::=
(**exhaustive search** | **AD-HOC method**)

Now, let us assume that there are m tuning regions with parameters P_i ($i=1, 2, \dots m$), each needing to vary N_i parameters. In this case, the parameters are expressed as follows:

$$P = (V(P_1), V(P_2), \dots, V(P_m)),$$

where $V(P_i)$ expresses 1 of the N_i parameters of parameter P_i .

Exhaustive search:

!OAT\$ search Brute-force

In an exhaustive search, all combinations are investigated. In other words, under this

method all combinations of parameter P are searched.

Consequently, the number of parameter combinations is ΠN_i .

AD-HOC method:

!OAT\$ search AD-HOC

When the AD-HOC method is used, not all combinations of parameter P are searched.

The search starts with a given parameter with a set initial value, a varied P_m , and the optimum parameter found and set. Next, it is P_{m-1} , then the optimum parameter is found and set, and so on. The algorithm then repeats the process until P_1 .

Consequently, the number of parameter combinations is ΣN_i .

Action When Different Search Methods are Specified for Different Nested Specifiers

In general, the search begins from the innermost AT region, and is made to match the outermost search method. However, if the inner method is AD-HOC, and the outer method is exhaustive, it will be treated as if the parameters of the AD-HOC specified AT regions are constant values.

Sample Program 10:

How is parameter searching carried out for the following nested processes?

```
!OAT$ static variable (BL) region start
!OAT$ name ABlockRoutine
!OAT$ varied 1 from 16
do iter=1, n, BL
...
!OAT$ static unroll (i,j) region start
!OAT$ name Kernel1
!OAT$ varied (i,j) from 1 to 32
  do i=1+iter, n
    do j=1+iter, n
      do k=1+iter, n
        .....
      enddo
    enddo
  enddo
!OAT$ static unroll (i,j) region end
....
!OAT$ static unroll (i,j) region start
!OAT$ name Kernel2
!OAT$ varied (l,m) from 1 to 32
  do l=1+iter, n
    do m=1+iter, n
      do p=1+iter, n
        .....
      enddo
    enddo
  enddo
!OAT$ static unroll (l,m) region end
...
enddo
!OAT$ static variable (BL) region end
```


Here, we assume that the parameter ordering is (BL, (i,j),(l,m)).

In the case of the above example, an exhaustive search is performed for all tuning regions: AblockRoutine, Kernel1, and Kernel2. Here, the parameter search proceeds as follows:

(1,(1,1),(1,1)), (1,(1,1),(1,2)), ..., (1,(1,1),(1,32)),
(1,(1,1),(2,1)), (1,(1,1),(2,2)), ..., (1,(1,1),(2,32)),
...,
(1,(1,2),(1,1)), (1,(1,2),(1,2)), ..., (1,(1,2),(1,32)), ...,

Thus, there are $16 \times 32 \times 32 \times 32 \times 32 = 1,677,216$ searches.

Let us assume that in the above example, the method for all tuning regions (AblockRoutine, Kernel1, and Kernel2) is AD-HOC. In this case, the search will be as follows:

(1,(1,1),(1,1)), (1,(1,1),(1,2)), ..., (1,(1,1),(1,32)): Fastest parameter is determined (e.g. 8)
(1,(1,1),(1,8)), (1,(1,1),(2,8)), ..., (1,(1,1),(32,8)): Fastest parameter is determined (e.g. 4)
(1,(1,1),(4,8)), (1,(1,2),(4,8)), ..., (1,(1,32),(4,8)): Fastest parameter is determined (e.g. 5)
(1,(1,5),(4,8)), (1,(2,5),(4,8)), ...

In other words, there are $16 + 32 + 32 + 32 + 32 = 144$ parameter searches.

Let us assume that in the above example, the method for tuning region AblockRoutine is exhaustive search, and that for tuning regions Kernel1 and Kernel2, it is AD-HOC. In this case, the search will be as follows:

(1,(1,1),(1,1)), (1,(1,1),(1,2)), ..., (1,(1,1),(1,32)): Fastest parameter is determined (e.g. 8)
(1,(1,1),(1,8)), (1,(1,1),(2,8)), ..., (1,(1,1),(32,8)): Fastest parameter is determined (e.g. 4)
(1,(1,1),(4,8)), (1,(1,2),(4,8)), ..., (1,(1,32),(4,8)): Fastest parameter is determined (e.g. 5)
(1,(1,5),(4,8)), (1,(2,5),(4,8)), ..., (1,(32,5),(4,8)): Fastest parameter is determined (e.g. 6)
(1,(6,5),(4,8)), (2,(6,5),(4,8)), ...,

In other words, there are $16 + 32 + 32 + 32 + 32 = 144$ parameter searches.

Now let us assume that in the above example, the method for tuning region AblockRoutine is AD-HOC, and that for tuning regions Kernel1 and Kernel2, it is

exhaustive search. In this case, the search will be as follows:

(1,(1,1),(1,1)),(1,(1,1),(1,2)),..., (1,(1,1),(1,32)) ,
(1,(1,1),(2,1)),(1,(1,1),(2,2)),..., (1,(1,1),(2,32)) ,
...
(1,(1,1),(32,1)),(1,(1,1),(32,2)),..., (1,(1,1),(32,32)): Fastest parameter is determined (e.g.
(3,9))
(1,(1,1),(3,9)),(1,(1,2),(3,9)),..., (1,(1,32),(3,9)),
(1,(2,1),(3,9)),(1,(2,2),(3,9)),..., (1,(2,32),(3,9)),
...
(1,(32,1),(3,9)),(1,(32,2),(3,9)),..., (1,(32,32),(3,9)): Fastest parameter is determined (e.g.
(2,8))
(1,(2,8),(3,9)),(2,(2,8),(3,9)),..., (16,(2,8),(3,9)): Fastest parameter is determined (e.g. 6)

In other words, there are $16+32*32+32*32 = 2,064$ parameter searches.

Note that if no search method is specified, the default methods are as follows.

<p>Feature: Default search method define: None (no need for search) variable: Exhaustive search select: AD-HOC search unroll: Exhaustive search</p>

Conclusion

This user's guide has described the specifications and usage of the *ppOpen-AT* auto-tuning processing directives for auto-tuning software developers.

The auto-tuning directives encoded in *ppOpen-AT* could be called the unique knowledge of software developers. Consequently, encoding the craftsman-like knowledge of software engineers using *ppOpen-AT* will pass on this knowledge to other engineers via the source code.

In the past, this knowledge had to be obtained individually, and was not something that could be published for others. That is what turned the expertise and knowledge of individuals into a kind of obscure craftsman-like technical knowledge.

In light of this background, the author believes that the clear encoding of knowledge in the source code using *ppOpen-AT* will have a major ripple effect.